

1-data-acquisition

April 8, 2024

```
[1]: # Aim : To perform Data Acquisition of given data set using Pandas
```

```
[2]: # Name: Samiksha Badhe  
# Class: 3rd Year  
# Sec: B  
# Roll No. : 05
```

```
[3]: import os  
import numpy as np  
import pandas as pd  
from sklearn.datasets import load_iris
```

```
[4]: data=pd.read_csv("C:\\Users\\hp\\Desktop\\IRIS.csv")
```

```
[5]: data.head()
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
[6]: data.tail()
```

	sepal_length	sepal_width	petal_length	petal_width	species
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 150 entries, 0 to 149  
Data columns (total 5 columns):  
#   Column          Non-Null Count  Dtype
```

```

---  -----
0   sepal_length  150 non-null    float64
1   sepal_width   150 non-null    float64
2   petal_length  150 non-null    float64
3   petal_width   150 non-null    float64
4   species       150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB

```

```
[8]: data.describe()
```

```

[8]:      sepal_length  sepal_width  petal_length  petal_width
count    150.000000    150.000000    150.000000    150.000000
mean         5.843333         3.054000         3.758667         1.198667
std          0.828066         0.433594         1.764420         0.763161
min          4.300000         2.000000         1.000000         0.100000
25%          5.100000         2.800000         1.600000         0.300000
50%          5.800000         3.000000         4.350000         1.300000
75%          6.400000         3.300000         5.100000         1.800000
max          7.900000         4.400000         6.900000         2.500000

```

```
[9]: data.ndim
```

```
[9]: 2
```

```
[10]: data.shape
```

```
[10]: (150, 5)
```

```
[11]: data.size
```

```
[11]: 750
```

```
[12]: data.isnull()
```

```

[12]:      sepal_length  sepal_width  petal_length  petal_width  species
0           False          False          False          False    False
1           False          False          False          False    False
2           False          False          False          False    False
3           False          False          False          False    False
4           False          False          False          False    False
..           ...           ...           ...           ...
145          False          False          False          False    False
146          False          False          False          False    False
147          False          False          False          False    False
148          False          False          False          False    False
149          False          False          False          False    False

```

[150 rows x 5 columns]

```
[13]: data.isnull().sum()
```

```
[13]: sepal_length    0
      sepal_width    0
      petal_length   0
      petal_width    0
      species        0
      dtype: int64
```

```
[14]: df = pd.DataFrame(data)

      # Count unique values in 'Column1'
      value_counts = df['species'].value_counts()
      print(value_counts)
```


```
species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
Name: count, dtype: int64
```

```
[15]: from matplotlib import pyplot as plt
```

```
[16]: iris = load_iris()
      target_names = iris.target_names
      y = iris.target

      # Count occurrences of each class
      class_counts = np.bincount(y)

      # Plot the bar chart
      plt.bar(target_names, class_counts, color=['blue', 'green', 'red'])
      plt.xlabel('Class')
      plt.ylabel('Count')
      plt.title('Distribution of Iris Classes')
      plt.show()
```



[]:

2-linear-regression

April 8, 2024

```
[1]: # Aim: To perform Simple Linear Regression and find out the Coefficient of it.
```

```
[2]: # Name: Samiksha Badhe  
# Class: 3rd Year  
# Section : B  
# Roll no: 05
```

```
[3]: import os
```

```
[4]: import pandas as pd
```

```
[5]: import numpy as np # linear algebra  
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
[6]: from sklearn.linear_model import LogisticRegression # for Logistic Regression  
      ↪algorithm  
from sklearn.model_selection import train_test_split  
from sklearn.datasets import load_iris  
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import accuracy_score  
from sklearn.metrics import mean_squared_error, r2_score
```

```
[7]: iris=load_iris()  
X = iris.data # Features  
y = iris.target  
dir(iris)
```

```
[7]: ['DESCR',  
      'data',  
      'data_module',  
      'feature_names',  
      'filename',  
      'frame',  
      'target',  
      'target_names']
```

```
[8]: os.getcwd()
```

```
[8]: 'C:\\Users\\hp\\Desktop\\BDA practicals(ET-2)'
```

```
[9]: df=pd.read_csv("C://Users//hp//Desktop//IRIS.csv")
```

```
[10]: df.head()
```

```
[10]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
[11]: df.tail()
```

```
[11]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

```
[12]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
[13]: df.describe()
```

```
[13]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000

50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
[14]: df.isnull()
```

```
[14]:      sepal_length  sepal_width  petal_length  petal_width  species
0             False           False           False           False    False
1             False           False           False           False    False
2             False           False           False           False    False
3             False           False           False           False    False
4             False           False           False           False    False
..            ...             ...             ...             ...      ...
145           False           False           False           False    False
146           False           False           False           False    False
147           False           False           False           False    False
148           False           False           False           False    False
149           False           False           False           False    False

[150 rows x 5 columns]
```

```
[15]: df.isna().sum()
```

```
[15]: sepal_length    0
      sepal_width    0
      petal_length    0
      petal_width    0
      species        0
      dtype: int64
```

```
[16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=40)
```

```
[17]: model = LinearRegression()

# Train the model using the training sets
model.fit(X_train, y_train)

# Make predictions using the testing set
y_pred = model.predict(X_test)

# Coefficients
print('Coefficients:', model.coef_)
```

```
Coefficients: [-0.1502982  -0.04339123  0.25345042  0.58205165]
```

```
[18]: mse = mean_squared_error(y_test, y_pred)
print('Mean squared error: %.2f' % mse)

# Calculate coefficient of determination (R^2 score)
r2 = r2_score(y_test, y_pred)
print('Coefficient of determination (R^2 score): %.2f' % r2)
```

Mean squared error: 0.04

Coefficient of determination (R^2 score): 0.94

```
[ ]:
```


3-logistic-regression

April 8, 2024

```
[1]: # Aim: To perform and find the accuracy of Logistic Regression
```

```
[2]: # Name: Samiksha Badhe  
# Class: 3rd Year  
# Section : B  
# Roll no: 05
```

```
[3]: import pandas as pd
```

```
[4]: from sklearn.model_selection import train_test_split
```

```
[5]: from sklearn.linear_model import LogisticRegression
```

```
[6]: from sklearn.metrics import accuracy_score
```

```
[7]: df=pd.read_csv('C:\\Users\\hp\\Desktop\\CHD_preprocessed.csv')
```

```
[8]: df.head()
```

```
[8]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	\
0	1	39	1	0	0.0	0.0	0	
1	0	46	0	0	0.0	0.0	0	
2	1	48	0	1	20.0	0.0	0	
3	0	61	1	1	30.0	0.0	0	
4	0	46	1	1	23.0	0.0	0	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0	0	0	195.0	106.0	70.0	26.97	80.0	77.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	76.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	70.0	
3	1	0	225.0	150.0	95.0	28.58	65.0	103.0	
4	0	0	285.0	130.0	84.0	23.10	85.0	85.0	

	TenYearCHD
0	0
1	0
2	0

```
3          1
4          0
```

```
[9]: df.tail()
```

```
[9]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
4128     1   50         0           1         1.0     0.0
4129     1   51         1           1        43.0     0.0
4130     0   48         0           1        20.0     0.0
4131     0   44         0           1        15.0     0.0
4132     0   52         0           0         0.0     0.0

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
4128                0            1         0    313.0  179.0   92.0  25.97
4129                0            0         0    207.0  126.5   80.0  19.71
4130                0            0         0    248.0  131.0   72.0  22.00
4131                0            0         0    210.0  126.5   87.0  19.16
4132                0            0         0    269.0  133.5   83.0  21.47

      heartRate  glucose  TenYearCHD
4128      66.0    86.0           1
4129      65.0    68.0           0
4130      84.0    86.0           0
4131      86.0    82.0           0
4132      80.0   107.0           0
```

```
[10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4133 entries, 0 to 4132
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4133 non-null  int64
1   age                   4133 non-null  int64
2   education             4133 non-null  int64
3   currentSmoker         4133 non-null  int64
4   cigsPerDay            4133 non-null  float64
5   BPMeds                4133 non-null  float64
6   prevalentStroke       4133 non-null  int64
7   prevalentHyp          4133 non-null  int64
8   diabetes              4133 non-null  int64
9   totChol               4133 non-null  float64
10  sysBP                 4133 non-null  float64
11  diaBP                 4133 non-null  float64
12  BMI                   4133 non-null  float64
13  heartRate             4133 non-null  float64
```

```

14 glucose          4133 non-null   float64
15 TenYearCHD       4133 non-null   int64
dtypes: float64(8), int64(8)
memory usage: 516.8 KB

```

```
[11]: df.describe()
```

```

[11]:
count      male      age  education  currentSmoker  cigsPerDay \
mean      0.427293  49.557222   0.280668      0.494798      9.101621
std       0.494745   8.561628   0.449380      0.500033     11.918440
min       0.000000  32.000000   0.000000      0.000000      0.000000
25%       0.000000  42.000000   0.000000      0.000000      0.000000
50%       0.000000  49.000000   0.000000      0.000000      0.000000
75%       1.000000  56.000000   1.000000      1.000000     20.000000
max       1.000000  70.000000   1.000000      1.000000     70.000000

count      BPMeds  prevalentStroke  prevalentHyp      diabetes      totChol \
mean      0.034358      0.006049      0.311154      0.025647     236.664408
std       0.182168      0.077548      0.463022      0.158100     43.909188
min       0.000000      0.000000      0.000000      0.000000     107.000000
25%       0.000000      0.000000      0.000000      0.000000     206.000000
50%       0.000000      0.000000      0.000000      0.000000     234.000000
75%       0.000000      0.000000      1.000000      0.000000     262.000000
max       1.000000      1.000000      1.000000      1.000000     600.000000

count      sysBP      diaBP      BMI      heartRate      glucose \
mean     132.367046     82.872248     25.778571     75.925236     81.946528
std       22.080332     11.952654      4.074360     12.049188     22.860954
min       83.500000     48.000000     15.540000     44.000000     40.000000
25%      117.000000     75.000000     23.060000     68.000000     72.000000
50%      128.000000     82.000000     25.380000     75.000000     80.000000
75%      144.000000     89.500000     27.990000     83.000000     85.000000
max      295.000000    142.500000     56.800000    143.000000    394.000000

count      TenYearCHD
mean      0.151948
std       0.359014
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000
max       1.000000

```

```
[12]: df.isnull()
```

```
[12]:      male    age  education  currentSmoker  cigsPerDay  BPMeds  \
0      False  False      False      False      False    False
1      False  False      False      False      False    False
2      False  False      False      False      False    False
3      False  False      False      False      False    False
4      False  False      False      False      False    False
...
4128   False  False      False      False      False    False
4129   False  False      False      False      False    False
4130   False  False      False      False      False    False
4131   False  False      False      False      False    False
4132   False  False      False      False      False    False

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
0              False      False    False    False  False  False  False
1              False      False    False    False  False  False  False
2              False      False    False    False  False  False  False
3              False      False    False    False  False  False  False
4              False      False    False    False  False  False  False
...
4128           False      False    False    False  False  False  False
4129           False      False    False    False  False  False  False
4130           False      False    False    False  False  False  False
4131           False      False    False    False  False  False  False
4132           False      False    False    False  False  False  False

      heartRate  glucose  TenYearCHD
0           False    False      False
1           False    False      False
2           False    False      False
3           False    False      False
4           False    False      False
...
4128         False    False      False
4129         False    False      False
4130         False    False      False
4131         False    False      False
4132         False    False      False
```

```
[4133 rows x 16 columns]
```

```
[13]: df.isna().sum()
```

```
[13]: male          0
      age          0
```

```

education      0
currentSmoker  0
cigsPerDay     0
BPMeds         0
prevalentStroke 0
prevalentHyp   0
diabetes       0
totChol        0
sysBP          0
diaBP          0
BMI            0
heartRate      0
glucose        0
TenYearCHD     0
dtype: int64

```

```

[14]: x=df.drop("TenYearCHD",axis=1)
      y=df['TenYearCHD']

```

```

[15]: x

```

```

[15]:
   male  age  education  currentSmoker  cigsPerDay  BPMeds  \
0      1   39         1             0         0.0     0.0
1      0   46         0             0         0.0     0.0
2      1   48         0             1        20.0     0.0
3      0   61         1             1        30.0     0.0
4      0   46         1             1        23.0     0.0
...    ...  ...      ...             ...         ...
4128    1   50         0             1         1.0     0.0
4129    1   51         1             1        43.0     0.0
4130    0   48         0             1        20.0     0.0
4131    0   44         0             1        15.0     0.0
4132    0   52         0             0         0.0     0.0

   prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
0                0             0         0    195.0   106.0   70.0  26.97
1                0             0         0    250.0   121.0   81.0  28.73
2                0             0         0    245.0   127.5   80.0  25.34
3                0             1         0    225.0   150.0   95.0  28.58
4                0             0         0    285.0   130.0   84.0  23.10
...              ...           ...         ...         ...         ...
4128              0             1         0    313.0   179.0   92.0  25.97
4129              0             0         0    207.0   126.5   80.0  19.71
4130              0             0         0    248.0   131.0   72.0  22.00
4131              0             0         0    210.0   126.5   87.0  19.16
4132              0             0         0    269.0   133.5   83.0  21.47

```

	heartRate	glucose
0	80.0	77.0
1	95.0	76.0
2	75.0	70.0
3	65.0	103.0
4	85.0	85.0
...
4128	66.0	86.0
4129	65.0	68.0
4130	84.0	86.0
4131	86.0	82.0
4132	80.0	107.0

[4133 rows x 15 columns]

[16]: y

```
[16]: 0      0
      1      0
      2      0
      3      1
      4      0

      ..
      4128    1
      4129    0
      4130    0
      4131    0
      4132    0
```

Name: TenYearCHD, Length: 4133, dtype: int64

[17]: x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.
↪2,random_state=42)

[18]: model = LogisticRegression(max_iter=1600)

[19]: model.fit(x_train,y_train)
model.score(x_train, y_train)

[19]: 0.8623714458560193

description-using-numpy-and-scipy

April 8, 2024

```
[1]: # Aim: To perform finding Stastical mean, median, mode, standard deviation,   
      ↪ Variance using Numpy and Scipy
```

```
[2]: # Name: Samiksha Badhe  
      # Class: 3rd Year  
      # Sec: B  
      # Roll No. : 05
```

```
[3]: import numpy as np  
      from scipy import stats
```

```
[4]: x=np.array([1,2,3,4,5,6,7,2,6,2,1,4,2,2,6])
```

```
[5]: x
```

```
[5]: array([1, 2, 3, 4, 5, 6, 7, 2, 6, 2, 1, 4, 2, 2, 6])
```

```
[6]: print(np.mean(x))
```

```
3.533333333333333
```

```
[7]: print(np.median(x))
```

```
3.0
```

```
[8]: print(stats.mode(x))
```

```
ModeResult(mode=2, count=5)
```

```
[9]: from scipy import stats
```

```
[10]: print(stats.mode(x))
```

```
ModeResult(mode=2, count=5)
```

```
[11]: print(np.std(x))
```

```
1.9618585292749546
```

```
[12]: print(np.var(x))
```

3.8488888888888884

```
[13]: import numpy as np
x=np.array([1,100,200,300,4000,5000])
y=np.array([2,4,6,8,10])
```

```
[14]: print(np.std(x))
```

2072.711623024829

```
[15]: print(np.std(y))
```

2.8284271247461903


```
[16]: print(np.var(x))
```

4296133.472222221


```
[17]: print(np.var(y))
```

8.0

```
[18]: from matplotlib import pyplot as plt
plt.hist(x)
plt.show()
```

```
[19]: from matplotlib import pyplot as plt  
plt.hist(y)  
plt.show()
```



```
[20]: from statsmodels.stats.weightstats import ztest as ztest
      #enter IQ levels for 20 patients
      data = [88, 92, 94, 94, 96, 97, 97, 97, 99, 99,
              105, 109, 109, 109, 110, 112, 112, 113, 114, 115]
      #perform one sample z-test
      ztest(data)
      (1.5976240527147705, 0.1101266701438426)
```

```
[20]: (1.5976240527147705, 0.1101266701438426)
```

5-f-test-anova

April 8, 2024

1 F-Test

```
[1]: # Aim : To perform hypothesis testing using ANOVA (F-TEST) One-Way
      ↪ F-Test(Anova).
```

```
[2]: # Name: Samiksha Badhe
      # Class: 3rd Year
      # Sec: B
      # Roll No. : 05
```

```
[3]: ages=[10,20,35,50,28,40,55,18,16,55,30,25,43,18,30,28,14,24,16,17,32,35,26,27,65,18,43,23,21,2
```

```
[4]: len(ages)
```

```
[4]: 56
```

```
[5]: import numpy as np
```

```
[6]: sample_size=10
      age_sample=np.random.choice(ages,sample_size)
```

```
[7]: import scipy.stats
      import numpy as np
```

```
[8]: data1 = [0.0842, 0.0368, 0.0847, 0.0935, 0.0376, 0.0963, 0.0684,
              0.0758, 0.0854, 0.0855]
      data2 = [0.0785, 0.0845, 0.0758, 0.0853, 0.0946, 0.0785, 0.0853,
              0.0685]
      data3 = [0.0864, 0.2522, 0.0894, 0.2724, 0.0853, 0.1367, 0.853]
```

```
[9]: # Performing the F-Test
      f_test, p_val = scipy.stats.f_oneway(data1, data2, data3)
      print("p-value is: ", p_val)
```

p-value is: 0.04043792126789144

```
[10]: # taking the threshold value as 0.05 or 5%
      if p_val < 0.05:
          print(" We can reject the null hypothesis")
      else:
          print("We can accept the null hypothesis")
```

We can reject the null hypothesis

```
[11]: variance1 = np.var(data1)
```

```
[12]: print(variance1)
```

0.00040949560000000005

```
[13]: variance2 = np.var(data2)
```

```
[14]: print(variance2)
```

5.3606874999999995e-05

```
[15]: variance3 = np.var(data3)
```

```
[16]: print(variance3)
```

0.06522053346938775

```
[ ]:
```

6-t-test

April 8, 2024

1 T Test

```
[1]: # Aim : To perform hypothesis testing using T test.
```

```
[2]: # Name: Samiksha Badhe  
# Class: 3rd Year  
# Sec: B  
# Roll No. : 05
```

T Test A t-test is a type of inferential statistic which is used to determine if there is a significant difference between the means of two groups which may be related in certain features

```
[3]: ages=[10,20,35,50,28,40,55,18,16,55,30,25,43,18,30,28,14,24,16,17,32,35,26,27,65,18,43,23,21,2
```

```
[4]: len(ages)
```

```
[4]: 32
```

```
[5]: import numpy as np  
ages_mean=np.mean(ages)  
print(ages_mean)
```

```
30.34375
```

```
[6]: sample_size=10  
age_sample=np.random.choice(ages,sample_size)
```

```
[7]: age_sample
```

```
[7]: array([18, 14, 70, 21, 50, 18, 18, 18, 50, 55])
```

```
[8]: from scipy.stats import ttest_1samp
```

```
[9]: ttest,p_value=ttest_1samp(age_sample,30)
```

```
[10]: print(p_value)
```

```
0.6357349574999751
```

```
[11]: if p_value < 0.05:    # alpha value is 0.05 or 5%
      print(" we are rejecting null hypothesis")
      else:
      print("we are accepting null hypothesis")
```

we are accepting null hypothesis

7-z-test

April 8, 2024

1 Z Test

```
[1]: # Aim : To perform hypothesis testing using Z test.
```

```
[2]: # Name: Samiksha Badhe  
# Class: 3rd Year  
# Sec: B  
# Roll No. : 05
```

```
[3]: ages=[10,20,35,50,28,40,55,18,16,55,30,25,43,18,30,28,14,24,16,17,32,35,26,27,65,18,43,23,21,2
```

```
[4]: len(ages)
```

```
[4]: 32
```

```
[5]: import numpy as np  
ages_mean=np.mean(ages)  
print(ages_mean)
```

```
30.34375
```

```
[6]: ## Lets take sample  
  
sample_size=31  
age_sample=np.random.choice(ages,sample_size)
```

```
[7]: age_sample
```

```
[7]: array([24, 24, 55, 16, 65, 55, 35, 14, 65, 20, 28, 43, 10, 10, 30, 30, 55,  
        32, 20, 17, 23, 32, 17, 70, 27, 16, 16, 18, 20, 19, 35])
```

```
[8]: # from scipy.stats import ztest_1samp
```

```
[9]: from statsmodels.stats import weightstats as stests  
  
# Perform one-sample z-test  
ztest, p_value = stests.ztest(age_sample)
```

```
# Print the results
print("ztest", ztest)
print("P-value:", p_value)
```

```
ztest 9.851701971870249
P-value: 6.73929402110435e-23
```

```
[10]: if p_value < 0.05:    # alpha value is 0.05 or 5%    (Level of significance)
      print(" we are rejecting null hypothesis")
      else:
      print("we are accepting null hypothesis")
```

```
we are rejecting null hypothesis
```

```
[ ]:
```


8-knn-classifier

April 8, 2024

```
[1]: # Aim: To perform and find the accuracy of K-Nearest Neighbors Algorithm i.e. ↵  
      ↵KNN Classifier
```

```
[2]: # Name: Samiksha Badhe  
      # Class: 3rd Year  
      # Sec: B  
      # Roll No. : 05
```

```
[3]: import pandas as pd  
      import os  
      import matplotlib.pyplot as plt  
      import numpy as np  
      import seaborn as sns  
      from sklearn.model_selection import train_test_split  
      import warnings  
      warnings.filterwarnings('ignore')
```

```
[4]: df=pd.read_csv('C:\\Users\\hp\\Desktop\\CHD_preprocessed.csv')
```

```
[5]: df.head()
```

```
[5]:   male  age  education  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \  
0      1   39          1              0         0.0     0.0              0  
1      0   46          0              0         0.0     0.0              0  
2      1   48          0              1        20.0     0.0              0  
3      0   61          1              1        30.0     0.0              0  
4      0   46          1              1        23.0     0.0              0  
  
      prevalentHyp  diabetes  totChol  sysBP  diaBP   BMI  heartRate  glucose  \  
0                0         0    195.0  106.0   70.0  26.97     80.0    77.0  
1                0         0    250.0  121.0   81.0  28.73     95.0    76.0  
2                0         0    245.0  127.5   80.0  25.34     75.0    70.0  
3                1         0    225.0  150.0   95.0  28.58     65.0   103.0  
4                0         0    285.0  130.0   84.0  23.10     85.0    85.0  
  
      TenYearCHD  
0                0
```

```

1      0
2      0
3      1
4      0

```

```
[6]: df.tail()
```

```

[6]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
4128     1   50         0             1         1.0     0.0
4129     1   51         1             1        43.0     0.0
4130     0   48         0             1        20.0     0.0
4131     0   44         0             1        15.0     0.0
4132     0   52         0             0         0.0     0.0

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
4128                0             1         0   313.0  179.0   92.0  25.97
4129                0             0         0   207.0  126.5   80.0  19.71
4130                0             0         0   248.0  131.0   72.0  22.00
4131                0             0         0   210.0  126.5   87.0  19.16
4132                0             0         0   269.0  133.5   83.0  21.47

      heartRate  glucose  TenYearCHD
4128      66.0    86.0           1
4129      65.0    68.0           0
4130      84.0    86.0           0
4131      86.0    82.0           0
4132      80.0   107.0           0

```

```
[7]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4133 entries, 0 to 4132
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4133 non-null  int64
1   age                   4133 non-null  int64
2   education              4133 non-null  int64
3   currentSmoker          4133 non-null  int64
4   cigsPerDay              4133 non-null  float64
5   BPMeds                 4133 non-null  float64
6   prevalentStroke         4133 non-null  int64
7   prevalentHyp            4133 non-null  int64
8   diabetes                4133 non-null  int64
9   totChol                 4133 non-null  float64
10  sysBP                  4133 non-null  float64
11  diaBP                  4133 non-null  float64

```

```

12 BMI                4133 non-null    float64
13 heartRate          4133 non-null    float64
14 glucose            4133 non-null    float64
15 TenYearCHD         4133 non-null    int64
dtypes: float64(8), int64(8)
memory usage: 516.8 KB

```

```
[8]: df.describe()
```

```

[8]:
count      male      age  education  currentSmoker  cigsPerDay  \
count  4133.000000  4133.000000  4133.000000    4133.000000  4133.000000
mean      0.427293   49.557222    0.280668      0.494798     9.101621
std       0.494745    8.561628    0.449380      0.500033    11.918440
min       0.000000   32.000000    0.000000      0.000000     0.000000
25%       0.000000   42.000000    0.000000      0.000000     0.000000
50%       0.000000   49.000000    0.000000      0.000000     0.000000
75%       1.000000   56.000000    1.000000      1.000000    20.000000
max       1.000000   70.000000    1.000000      1.000000    70.000000

count      BPMeds  prevalentStroke  prevalentHyp  diabetes  totChol  \
count  4133.000000    4133.000000    4133.000000  4133.000000  4133.000000
mean      0.034358      0.006049      0.311154      0.025647   236.664408
std       0.182168      0.077548      0.463022      0.158100   43.909188
min       0.000000      0.000000      0.000000      0.000000  107.000000
25%       0.000000      0.000000      0.000000      0.000000  206.000000
50%       0.000000      0.000000      0.000000      0.000000  234.000000
75%       0.000000      0.000000      1.000000      0.000000  262.000000
max       1.000000      1.000000      1.000000      1.000000  600.000000

count      sysBP      diaBP      BMI      heartRate      glucose  \
count  4133.000000  4133.000000  4133.000000  4133.000000  4133.000000
mean    132.367046    82.872248    25.778571    75.925236    81.946528
std     22.080332    11.952654     4.074360    12.049188    22.860954
min     83.500000    48.000000    15.540000    44.000000    40.000000
25%    117.000000    75.000000    23.060000    68.000000    72.000000
50%    128.000000    82.000000    25.380000    75.000000    80.000000
75%    144.000000    89.500000    27.990000    83.000000    85.000000
max    295.000000   142.500000    56.800000   143.000000   394.000000

count      TenYearCHD
count  4133.000000
mean      0.151948
std       0.359014
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000

```

```
max          1.000000
```

```
[9]: df.isna().sum()
```

```
[9]: male          0
age            0
education      0
currentSmoker  0
cigsPerDay     0
BPMeds         0
prevalentStroke 0
prevalentHyp   0
diabetes       0
totChol        0
sysBP          0
diaBP          0
BMI            0
heartRate      0
glucose        0
TenYearCHD     0
dtype: int64
```

```
[18]: # Splitting the dependent and independent variables
x = df.drop('TenYearCHD',axis=1)
y = df['TenYearCHD']
```

```
[19]: x #Checking the features
```

```
[19]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	\
0	1	39	1	0	0.0	0.0	
1	0	46	0	0	0.0	0.0	
2	1	48	0	1	20.0	0.0	
3	0	61	1	1	30.0	0.0	
4	0	46	1	1	23.0	0.0	
...	
4128	1	50	0	1	1.0	0.0	
4129	1	51	1	1	43.0	0.0	
4130	0	48	0	1	20.0	0.0	
4131	0	44	0	1	15.0	0.0	
4132	0	52	0	0	0.0	0.0	

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
0	0	0	0	195.0	106.0	70.0	26.97	
1	0	0	0	250.0	121.0	81.0	28.73	
2	0	0	0	245.0	127.5	80.0	25.34	
3	0	1	0	225.0	150.0	95.0	28.58	
4	0	0	0	285.0	130.0	84.0	23.10	

...
4128	0	1	0	313.0	179.0	92.0	25.97
4129	0	0	0	207.0	126.5	80.0	19.71
4130	0	0	0	248.0	131.0	72.0	22.00
4131	0	0	0	210.0	126.5	87.0	19.16
4132	0	0	0	269.0	133.5	83.0	21.47

	heartRate	glucose
0	80.0	77.0
1	95.0	76.0
2	75.0	70.0
3	65.0	103.0
4	85.0	85.0
...
4128	66.0	86.0
4129	65.0	68.0
4130	84.0	86.0
4131	86.0	82.0
4132	80.0	107.0

[4133 rows x 15 columns]

1 Train Test Split

```
[20]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
      ↪2,random_state=42)
```

```
[21]: y_train
```

```
[21]: 173      1
      1022    0
      3182    0
      331     1
      2222    0
      ..
      3444    0
      466     0
      3092    0
      3772    0
      860     0
      Name: TenYearCHD, Length: 3306, dtype: int64
```

```
[22]: y_test
```

```
[22]: 1864     0
      1210     0
```

```

1924    0
1752    0
1095    0
..
881     0
25      1
3256    0
2269    0
1074    0
Name: TenYearCHD, Length: 827, dtype: int64

```

```
[23]: x_train
```

```

[23]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
173      0   60           1             0         0.0     0.0
1022     1   42           1             1        20.0     0.0
3182     1   58           0             0         0.0     0.0
331      0   58           0             0         0.0     0.0
2222     1   39           1             0         0.0     0.0
...    ...  ...  ...      ...      ...      ...
3444     0   49           0             0         0.0     0.0
466      1   50           0             0         0.0     0.0
3092     0   36           0             0         0.0     0.0
3772     0   64           0             0         0.0     0.0
860      0   47           0             0         0.0     0.0

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
173                  0             1         0    325.0   182.0  106.0  27.61
1022                 0             0         0    270.0   112.0   77.0  24.77
3182                 0             0         0    225.0   105.5   74.0  25.68
331                  0             1         0    200.0   158.0  101.0  23.06
2222                 0             1         0    208.0   146.0   92.0  25.91
...                ...           ...      ...      ...      ...
3444                 0             1         0    233.0   149.0   91.5  26.03
466                  0             1         0    219.0   145.0  100.0  26.26
3092                 0             0         0    209.0   107.0   73.5  21.59
3772                 0             1         0    279.0   172.0   87.0  24.01
860                  0             0         0    232.0   113.5   73.0  28.78

      heartRate  glucose
173          80.0    77.0
1022         73.0    85.0
3182         50.0    93.0
331          85.0    77.0
2222         69.0    74.0
...          ...     ...
3444         68.0    82.0

```

466	78.0	108.0
3092	75.0	73.0
3772	80.0	70.0
860	75.0	77.0

[3306 rows x 15 columns]

[24]: x_test

```
[24]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	\
1864	1	40	0	1	5.0	0.0	
1210	0	50	0	1	10.0	0.0	
1924	0	64	1	0	0.0	0.0	
1752	0	55	0	0	0.0	1.0	
1095	0	46	1	0	0.0	0.0	
...	
881	0	44	0	1	1.0	0.0	
25	1	47	1	1	20.0	0.0	
3256	0	63	1	0	0.0	0.0	
2269	1	40	0	1	20.0	0.0	
1074	1	57	1	0	0.0	1.0	

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
1864	0	0	0	282.0	120.0	87.0	22.98	
1210	0	1	0	298.0	156.0	90.0	24.24	
1924	0	0	0	330.0	108.0	82.0	23.09	
1752	0	1	0	285.0	158.0	98.0	30.23	
1095	0	1	0	259.0	173.0	102.0	27.22	
...	
881	0	0	0	217.0	124.5	82.0	22.36	
25	0	0	0	294.0	102.0	68.0	24.18	
3256	0	0	0	297.0	133.5	92.0	25.09	
2269	0	0	0	193.0	122.0	78.0	28.40	
1074	0	1	0	195.0	162.0	108.0	32.65	

	heartRate	glucose
1864	60.0	82.0
1210	75.0	100.0
1924	85.0	80.0
1752	70.0	88.0
1095	85.0	75.0
...
881	87.0	68.0
25	62.0	66.0
3256	80.0	74.0
2269	70.0	93.0
1074	85.0	73.0

[827 rows x 15 columns]

```
[25]: from sklearn.neighbors import KNeighborsClassifier
      knn = KNeighborsClassifier(n_neighbors=5, p=2, metric='minkowski')
      knn.fit(x_train, y_train)
      acc = knn.score(x_test, y_test)*100
      print(acc)
```

81.62031438935912

[]:

9-k-means

April 8, 2024

```
[1]: #Aim: To perform and find the accuracy of K means algorithm
```

```
[ ]: # Name: Samiksha Badhe  
# Class: 3rd Year  
# Sec: B  
# Roll No. : 05
```

Running cells with 'c:\Users\ASUS-PC\AppData\Local\Microsoft\WindowsApps\python3.11.exe' requires the ipykernel package.

Run the following command to install 'ipykernel' into the Python environment.

Command: 'c:/Users/ASUS-PC/AppData/Local/Microsoft/WindowsApps/python3.11.exe -i -c "pip install ipykernel -U --user --force-reinstall"'

```
[2]: import numpy as np # linear algebra  
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)  
import matplotlib.pyplot as plt # for data visualization  
import seaborn as sns # for statistical data visualization  
%matplotlib inline  
from sklearn.cluster import KMeans  
from sklearn.metrics import adjusted_rand_score  
from sklearn.cluster import KMeans  
import warnings  
warnings.filterwarnings('ignore')
```

```
[3]: df=pd.read_csv('C:\\Users\\hp\\Desktop\\CHD_preprocessed.csv')
```

```
[4]: df.head()
```

```
[4]:   male  age  education  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \  
0     1   39          1              0          0.0    0.0              0  
1     0   46          0              0          0.0    0.0              0  
2     1   48          0              1         20.0    0.0              0  
3     0   61          1              1         30.0    0.0              0  
4     0   46          1              1         23.0    0.0              0
```

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0	0	0	195.0	106.0	70.0	26.97	80.0	77.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	76.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	70.0	
3	1	0	225.0	150.0	95.0	28.58	65.0	103.0	
4	0	0	285.0	130.0	84.0	23.10	85.0	85.0	

	TenYearCHD
0	0
1	0
2	0
3	1
4	0

```
[5]: df.tail()
```

```
[5]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
4128    1   50         0          1          1.0      0.0
4129    1   51         1          1         43.0      0.0
4130    0   48         0          1         20.0      0.0
4131    0   44         0          1         15.0      0.0
4132    0   52         0          0          0.0      0.0
```

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
4128	0	1	0	313.0	179.0	92.0	25.97	
4129	0	0	0	207.0	126.5	80.0	19.71	
4130	0	0	0	248.0	131.0	72.0	22.00	
4131	0	0	0	210.0	126.5	87.0	19.16	
4132	0	0	0	269.0	133.5	83.0	21.47	

	heartRate	glucose	TenYearCHD
4128	66.0	86.0	1
4129	65.0	68.0	0
4130	84.0	86.0	0
4131	86.0	82.0	0
4132	80.0	107.0	0

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4133 entries, 0 to 4132
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   male            4133 non-null   int64
1   age             4133 non-null   int64
```

```

2   education      4133 non-null   int64
3   currentSmoker  4133 non-null   int64
4   cigsPerDay     4133 non-null   float64
5   BPMeds         4133 non-null   float64
6   prevalentStroke 4133 non-null   int64
7   prevalentHyp   4133 non-null   int64
8   diabetes       4133 non-null   int64
9   totChol        4133 non-null   float64
10  sysBP          4133 non-null   float64
11  diaBP          4133 non-null   float64
12  BMI            4133 non-null   float64
13  heartRate      4133 non-null   float64
14  glucose        4133 non-null   float64
15  TenYearCHD     4133 non-null   int64
dtypes: float64(8), int64(8)
memory usage: 516.8 KB

```

```
[7]: df.size
```

```
[7]: 66128
```

```
[8]: df.shape
```

```
[8]: (4133, 16)
```

```
[9]: df.describe()
```

```

[9]:
count      male      age  education  currentSmoker  cigsPerDay  \
mean      0.427293  49.557222   0.280668      0.494798      9.101621
std       0.494745   8.561628   0.449380      0.500033     11.918440
min       0.000000  32.000000   0.000000      0.000000      0.000000
25%       0.000000  42.000000   0.000000      0.000000      0.000000
50%       0.000000  49.000000   0.000000      0.000000      0.000000
75%       1.000000  56.000000   1.000000      1.000000     20.000000
max       1.000000  70.000000   1.000000      1.000000     70.000000

count      BPMeds  prevalentStroke  prevalentHyp      diabetes      totChol  \
mean      0.034358      0.006049      0.311154      0.025647     236.664408
std       0.182168      0.077548      0.463022      0.158100     43.909188
min       0.000000      0.000000      0.000000      0.000000     107.000000
25%       0.000000      0.000000      0.000000      0.000000     206.000000
50%       0.000000      0.000000      0.000000      0.000000     234.000000
75%       0.000000      0.000000      1.000000      0.000000     262.000000
max       1.000000      1.000000      1.000000      1.000000     600.000000

```

	sysBP	diaBP	BMI	heartRate	glucose \
count	4133.000000	4133.000000	4133.000000	4133.000000	4133.000000
mean	132.367046	82.872248	25.778571	75.925236	81.946528
std	22.080332	11.952654	4.074360	12.049188	22.860954
min	83.500000	48.000000	15.540000	44.000000	40.000000
25%	117.000000	75.000000	23.060000	68.000000	72.000000
50%	128.000000	82.000000	25.380000	75.000000	80.000000
75%	144.000000	89.500000	27.990000	83.000000	85.000000
max	295.000000	142.500000	56.800000	143.000000	394.000000

	TenYearCHD
count	4133.000000
mean	0.151948
std	0.359014
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
[10]: X = df.drop(columns=['TenYearCHD'])
      kmeans = KMeans(n_clusters=2, random_state=0)
```

```
[11]: kmeans.fit(X)
      kmeans.cluster_centers_
      kmeans.inertia_
```

```
[11]: 9282994.90372527
```

```
[ ]:
```

10-naive-bayes-classifier

April 8, 2024

```
[1]: # Aim: To perform and find the accuracy of Naive bayes Classifier
```

```
[2]: # Name: Samiksha Badhe  
# Class: 3rd Year  
# Sec: B  
# Roll No. : 05
```

```
[3]: import pandas as pd  
import os  
import matplotlib.pyplot as plt  
import numpy as np  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
from sklearn.naive_bayes import GaussianNB  
import warnings  
warnings.filterwarnings('ignore')
```

```
[4]: df=pd.read_csv('C:\\Users\\hp\\Desktop\\CHD_preprocessed.csv')
```

```
[5]: df.head()
```

```
[5]:   male  age  education  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \  
0     1   39         1             0         0.0     0.0             0  
1     0   46         0             0         0.0     0.0             0  
2     1   48         0             1        20.0     0.0             0  
3     0   61         1             1        30.0     0.0             0  
4     0   46         1             1        23.0     0.0             0  
  
   prevalentHyp  diabetes  totChol  sysBP  diaBP   BMI  heartRate  glucose  \  
0              0         0   195.0  106.0   70.0  26.97      80.0    77.0  
1              0         0   250.0  121.0   81.0  28.73      95.0    76.0  
2              0         0   245.0  127.5   80.0  25.34      75.0    70.0  
3              1         0   225.0  150.0   95.0  28.58      65.0   103.0  
4              0         0   285.0  130.0   84.0  23.10      85.0    85.0  
  
   TenYearCHD  
0            0
```

```

1      0
2      0
3      1
4      0

```

```
[6]: df.tail()
```

```

[6]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
4128     1   50         0             1         1.0     0.0
4129     1   51         1             1        43.0     0.0
4130     0   48         0             1        20.0     0.0
4131     0   44         0             1        15.0     0.0
4132     0   52         0             0         0.0     0.0

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
4128                0             1         0   313.0  179.0   92.0  25.97
4129                0             0         0   207.0  126.5   80.0  19.71
4130                0             0         0   248.0  131.0   72.0  22.00
4131                0             0         0   210.0  126.5   87.0  19.16
4132                0             0         0   269.0  133.5   83.0  21.47

      heartRate  glucose  TenYearCHD
4128      66.0    86.0           1
4129      65.0    68.0           0
4130      84.0    86.0           0
4131      86.0    82.0           0
4132      80.0   107.0           0

```

```
[7]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4133 entries, 0 to 4132
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype
---  -
0   male                4133 non-null  int64
1   age                 4133 non-null  int64
2   education           4133 non-null  int64
3   currentSmoker       4133 non-null  int64
4   cigsPerDay          4133 non-null  float64
5   BPMeds              4133 non-null  float64
6   prevalentStroke     4133 non-null  int64
7   prevalentHyp        4133 non-null  int64
8   diabetes            4133 non-null  int64
9   totChol             4133 non-null  float64
10  sysBP               4133 non-null  float64
11  diaBP               4133 non-null  float64

```

```

12 BMI                4133 non-null    float64
13 heartRate          4133 non-null    float64
14 glucose            4133 non-null    float64
15 TenYearCHD         4133 non-null    int64
dtypes: float64(8), int64(8)
memory usage: 516.8 KB

```

```
[8]: df.describe()
```

```

[8]:
      count      male      age  education  currentSmoker  cigsPerDay  \
count  4133.000000  4133.000000  4133.000000    4133.000000  4133.000000
mean     0.427293   49.557222    0.280668     0.494798     9.101621
std     0.494745    8.561628    0.449380     0.500033    11.918440
min     0.000000   32.000000    0.000000     0.000000     0.000000
25%     0.000000   42.000000    0.000000     0.000000     0.000000
50%     0.000000   49.000000    0.000000     0.000000     0.000000
75%     1.000000   56.000000    1.000000     1.000000    20.000000
max     1.000000   70.000000    1.000000     1.000000    70.000000

      count  BPMeds  prevalentStroke  prevalentHyp  diabetes  totChol  \
count  4133.000000    4133.000000    4133.000000  4133.000000  4133.000000
mean     0.034358     0.006049     0.311154     0.025647   236.664408
std     0.182168     0.077548     0.463022     0.158100    43.909188
min     0.000000     0.000000     0.000000     0.000000   107.000000
25%     0.000000     0.000000     0.000000     0.000000   206.000000
50%     0.000000     0.000000     0.000000     0.000000   234.000000
75%     0.000000     0.000000     1.000000     0.000000   262.000000
max     1.000000     1.000000     1.000000     1.000000   600.000000

      count      sysBP      diaBP      BMI      heartRate      glucose  \
count  4133.000000  4133.000000  4133.000000  4133.000000  4133.000000
mean    132.367046   82.872248   25.778571   75.925236   81.946528
std     22.080332   11.952654    4.074360   12.049188   22.860954
min     83.500000   48.000000   15.540000   44.000000   40.000000
25%    117.000000   75.000000   23.060000   68.000000   72.000000
50%    128.000000   82.000000   25.380000   75.000000   80.000000
75%    144.000000   89.500000   27.990000   83.000000   85.000000
max    295.000000  142.500000   56.800000  143.000000  394.000000

      count  TenYearCHD
count  4133.000000
mean     0.151948
std     0.359014
min     0.000000
25%     0.000000
50%     0.000000
75%     0.000000

```

```
max      1.000000
```

```
[9]: df.size
```

```
[9]: 66128
```

```
[10]: df.shape
```

```
[10]: (4133, 16)
```

```
[11]: df.isna().sum()
```

```
[11]: male      0
age      0
education  0
currentSmoker  0
cigsPerDay  0
BPMeds     0
prevalentStroke  0
prevalentHyp  0
diabetes    0
totChol     0
sysBP       0
diaBP       0
BMI         0
heartRate   0
glucose     0
TenYearCHD  0
dtype: int64
```

```
[12]: x = df.drop("TenYearCHD",axis=1)
y = df['TenYearCHD']
```

```
[13]: x
```

```
[13]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	\
0	1	39	1	0	0.0	0.0	
1	0	46	0	0	0.0	0.0	
2	1	48	0	1	20.0	0.0	
3	0	61	1	1	30.0	0.0	
4	0	46	1	1	23.0	0.0	
...	
4128	1	50	0	1	1.0	0.0	
4129	1	51	1	1	43.0	0.0	
4130	0	48	0	1	20.0	0.0	
4131	0	44	0	1	15.0	0.0	
4132	0	52	0	0	0.0	0.0	

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
0	0	0	0	195.0	106.0	70.0	26.97	
1	0	0	0	250.0	121.0	81.0	28.73	
2	0	0	0	245.0	127.5	80.0	25.34	
3	0	1	0	225.0	150.0	95.0	28.58	
4	0	0	0	285.0	130.0	84.0	23.10	
...	
4128	0	1	0	313.0	179.0	92.0	25.97	
4129	0	0	0	207.0	126.5	80.0	19.71	
4130	0	0	0	248.0	131.0	72.0	22.00	
4131	0	0	0	210.0	126.5	87.0	19.16	
4132	0	0	0	269.0	133.5	83.0	21.47	

	heartRate	glucose
0	80.0	77.0
1	95.0	76.0
2	75.0	70.0
3	65.0	103.0
4	85.0	85.0
...
4128	66.0	86.0
4129	65.0	68.0
4130	84.0	86.0
4131	86.0	82.0
4132	80.0	107.0

[4133 rows x 15 columns]

[14]: y

```
[14]: 0      0
      1      0
      2      0
      3      1
      4      0
      ..
      4128    1
      4129    0
      4130    0
      4131    0
      4132    0
      Name: TenYearCHD, Length: 4133, dtype: int64
```

[15]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
↪2,random_state=42)

```
[16]: y_train
```

```
[16]: 173      1
      1022    0
      3182    0
      331     1
      2222    0
      ..
      3444    0
      466     0
      3092    0
      3772    0
      860     0
      Name: TenYearCHD, Length: 3306, dtype: int64
```

```
[17]: y_test
```

```
[17]: 1864     0
      1210     0
      1924     0
      1752     0
      1095     0
      ..
      881     0
      25      1
      3256     0
      2269     0
      1074     0
      Name: TenYearCHD, Length: 827, dtype: int64
```

```
[18]: nb_model = GaussianNB()
      nb_model.fit(x_train, y_train)
```

```
[18]: GaussianNB()
```

```
[19]: # Evaluate the model
      train_accuracy = nb_model.score(x_train, y_train)
      test_accuracy = nb_model.score(x_test, y_test)
```

```
[20]: print("Training Accuracy:", train_accuracy)
      print("Testing Accuracy:", test_accuracy)
```

```
Training Accuracy: 0.8236539624924379
Testing Accuracy: 0.8101571946795647
```

11-data-visualization

April 8, 2024

```
[1]: import numpy as np
      from matplotlib import pyplot as plt
```

```
[2]: # Name: Samiksha Badhe
      # Class: 3rd Year
      # Sec: B
      # Roll No. : 05
```

```
[3]: x=np.arange(1,11)
```

```
[4]: x
```

```
[4]: array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10])
```

```
[5]: y=x*2
```


```
[6]: x
```

```
[6]: array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10])
```


```
[7]: y
```

```
[7]: array([ 2,  4,  6,  8, 10, 12, 14, 16, 18, 20])
```


```
[8]: #Line chart
      plt.plot(x,y)
      plt.title("Line chart")
      plt.xlabel("X axis")
      plt.ylabel("Y axis")
      plt.show()
```




```
[9]: plt.bar(x,y)
plt.title("Bar chart")
plt.xlabel("X axis")
plt.ylabel("Y axis")
plt.show()
```




```
[10]: plt.scatter(x,y)
plt.title("Scatter plot")
plt.xlabel("X axis")
plt.ylabel("Y axis")
plt.show()
```



```
[11]: X=(1,9,2,8,3,7,4,7,5,6,8)
      Y=(9,7,5,3,1,2,4,6,8,0,1)
      plt.scatter(X,Y)
      plt.title("Scatter plot")
      plt.xlabel("X axis")
      plt.ylabel("Y axis")
      plt.show()
```




```
[12]: #Histogram  
  
H=[1,1,1,1,1,2,2,2,2,2,3,3,3,3,4,4,4,4,5,5,5,6,6,7,7,8,8,9,9,10,10,2,2,3,3,0,0,0,0,0]  
  
plt.hist(H,color="orange")  
plt.show()
```




```
[13]: # Box plot
n=[1,2,3,4,5,6,7,8,9]

plt.boxplot(n)
plt.show()
```

```
[14]: #pie chart
n=[1,2,3,4,5,6,7,8,9]
labels=['a','b','c','d','e','f','g','h','i']
plt.pie(n, labels=labels, autopct='%1.1f%%', startangle=140)
plt.show()
```



[]:

Practical No. 12

Aim : To study of Hadoop ecosystem in detail

Theory :

Apache Hadoop is an open source framework intended to make interaction with **big data** easier, However, for those who are not acquainted with this technology, one question arises that what is big data ? Big data is a term given to the data sets which can't be processed in an efficient manner with the help of traditional methodology such as RDBMS. Hadoop has made its place in the industries and companies that need to work on large data sets which are sensitive and needs efficient handling. Hadoop is a framework that enables processing of large data sets which reside in the form of clusters. Being a framework, Hadoop is made up of several modules that are supported by a large ecosystem of technologies.


Hadoop Ecosystem is a platform or a suite which provides various services to solve the big data problems. It includes Apache projects and various commercial tools and solutions. There are *four major elements of Hadoop* i.e. **HDFS, MapReduce, YARN, and Hadoop Common Utilities**. Most of the tools or solutions are used to supplement or support these major elements. All these tools work collectively to provide services such as absorption, analysis, storage and maintenance of data etc.

Following are the components that collectively form a Hadoop ecosystem:

Following are the components that collectively form a Hadoop ecosystem:

- **HDFS:** Hadoop Distributed File System
- **YARN:** Yet Another Resource Negotiator
- **MapReduce:** Programming based Data Processing
- **Spark:** In-Memory data processing
- **PIG, HIVE:** Query based processing of data services
- **HBase:** NoSQL Database


- **Mahout, Spark MLlib:** [Machine Learning](#) algorithm libraries
- **Solar, Lucene:** Searching and Indexing
- **Zookeeper:** Managing cluster
- **Oozie:** Job Scheduling



All these toolkits or components revolve around one term i.e. *Data*. That's the beauty of Hadoop that it revolves around data and hence making its synthesis easier.

HDFS:

- HDFS is the primary or major component of Hadoop ecosystem and is responsible for storing large data sets of structured or unstructured data across various nodes and thereby maintaining the metadata in the form of log files.
- HDFS consists of two core components i.e.
 - 1.Name node
 - 2.Data node
- Name Node is the prime node which contains metadata (data about data) requiring comparatively fewer resources than the data nodes that stores the actual data. These data nodes are commodity hardware in the distributed environment. Undoubtedly, making Hadoop cost effective
- HDFS maintains all the coordination between the clusters and hardware, thus working at the heart of the system.



YARN:

- Yet Another Resource Negotiator, as the name implies, YARN is the one who helps to manage the resources across the clusters. In short, it performs scheduling and resource allocation for the Hadoop System.
- Consists of three major components i.e.
 1. Resource Manager
 2. Nodes Manager
 3. Application Manager
- Resource manager has the privilege of allocating resources for the applications in a system whereas Node managers work on the allocation of resources such as CPU, memory, bandwidth per machine and later on acknowledges the resource manager. Application manager works as an interface between the resource manager and node manager and performs negotiations as per the requirement of the two.



MapReduce:

- By making the use of distributed and parallel algorithms, MapReduce makes it possible to carry over the processing's logic and helps to write applications which transform big data sets into a manageable one.
- MapReduce makes the use of two functions i.e. Map() and Reduce() whose task is:
 1. *Map()* performs sorting and filtering of data and thereby organizing them in the form of group. Map generates a key-value pair based result which is later on processed by the Reduce() method.
 2. *Reduce()*, as the name suggests does the summarization by aggregating the mapped data. In simple, Reduce() takes the output generated by Map() as input and combines those tuples into smaller set of tuples.

PIG:

Pig was basically developed by Yahoo which works on a pig Latin language, which is Query based language similar to SQL.

- It is a platform for structuring the data flow, processing and analyzing huge data sets.
- Pig does the work of executing commands and in the background, all the activities of MapReduce are taken care of. After the processing, pig stores the result in HDFS.

- Pig Latin language is specially designed for this framework which runs on Pig Runtime. Just the way Java runs on the [JVM](#).
- Pig helps to achieve ease of programming and optimization and hence is a major segment of the Hadoop Ecosystem.

HIVE:

- With the help of SQL methodology and interface, HIVE performs reading and writing of large data sets. However, its query language is called as HQL (Hive Query Language).
- It is highly scalable as it allows real-time processing and batch processing both. Also, all the SQL datatypes are supported by Hive thus, making the query processing easier.
- Similar to the Query Processing frameworks, HIVE too comes with two components: *JDBC Drivers* and *HIVE Command Line*.
- JDBC, along with ODBC drivers work on establishing the data storage permissions and connection whereas HIVE Command line helps in the processing of queries.

Mahout:

- Mahout, allows Machine Learnability to a system or application. [Machine Learning](#), as the name suggests helps the system to develop itself based on some patterns, user/environmental interaction or on the basis of algorithms.
- It provides various libraries or functionalities such as collaborative filtering, clustering, and classification which are nothing but concepts of Machine learning. It allows invoking algorithms as per our need with the help of its own libraries.

Apache Spark:

- It's a platform that handles all the process consumptive tasks like batch processing, interactive or iterative real-time processing, graph conversions, and visualization, etc.

- It consumes in memory resources hence, thus being faster than the prior in terms of optimization.
- Spark is best suited for real-time data whereas Hadoop is best suited for structured data or batch processing, hence both are used in most of the companies interchangeably.

Apache HBase:

- It's a NoSQL database which supports all kinds of data and thus capable of handling anything of Hadoop Database. It provides capabilities of Google's BigTable, thus able to work on Big Data sets effectively.
- At times where we need to search or retrieve the occurrences of something small in a huge database, the request must be processed within a short quick span of time. At such times, HBase comes handy as it gives us a tolerant way of storing limited data

Other Components: Apart from all of these, there are some other components too that carry out a huge task in order to make Hadoop capable of processing large datasets. They are as follows:

- **Solr, Lucene:** These are the two services that perform the task of searching and indexing with the help of some java libraries, especially Lucene is based on Java which allows spell check mechanism, as well. However, Lucene is driven by Solr.
- **Zookeeper:** There was a huge issue of management of coordination and synchronization among the resources or the components of Hadoop which resulted in inconsistency, often. Zookeeper overcame all the problems by performing synchronization, inter-component based communication, grouping, and maintenance.
- **Oozie:** Oozie simply performs the task of a scheduler, thus scheduling jobs and binding them together as a single unit. There is two kinds of jobs .i.e Oozie workflow and Oozie coordinator jobs. Oozie workflow is the jobs that need to be executed in a sequentially ordered manner whereas Oozie Coordinator jobs are those that are triggered when some data or external stimulus is given to it.

Conclusion : In this way , we learn about the Hadoop ecosystem.